

Boosting functional regression models with FDboost


Sarah Brockhaus & David Rügamer

In collaboration with Sonja Greven, Thorsten Hothorn, Andy Mayr,
Fabian Scheipl and Almond Stöcker

LMU Munich


July 24, 2017

Functional data: Growth curves



(Ramsay and Silverman, 2005)

Functional data: Spectrometric measures





(Brockhaus et al., 2015)

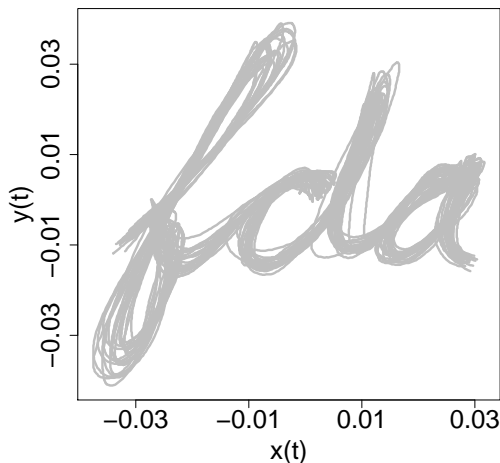
Functional data: Shapes



Functional data: Shapes




Functional data: Trajectories




(Ramsay and Silverman, 2005)




Functional data: Movement



Functional data: Brain scans






Introduction to functional data






- Observation units are functions; several measurement points for each observation unit

Introduction to functional data






- ▶ Observation units are functions; several measurement points for each observation unit
- ▶ Measurement points on regular or irregular grid

Introduction to functional data






- ▶ Observation units are functions; several measurement points for each observation unit
- ▶ Measurement points on regular or irregular grid
- ▶ Possibly arbitrary many measurements possible
→ smooth data generating function

Introduction to functional data



- ▶ Observation units are functions; several measurement points for each observation unit
- ▶ Measurement points on regular or irregular grid
- ▶ Possibly arbitrary many measurements possible
→ smooth data generating function
- ▶ Observations possibly with (measurement) error

Introduction to functional data



- ▶ Observation units are functions; several measurement points for each observation unit
- ▶ Measurement points on regular or irregular grid
- ▶ Possibly arbitrary many measurements possible
→ smooth data generating function
- ▶ Observations possibly with (measurement) error
- ▶ Difference: functional data \leftrightarrow time series

Outline

Functional data analysis in a nutshell

- Some basic statistics

- Overview

Regression with functional data

- Generic model

- Estimation by gradient boosting

- Other transformations of the conditional response distribution

- Implementation in FDboost

Case studies

- Functional response

- Scalar response and functional covariates

Summary and discussion

Basic statistics for functional data

Mean, Variance and Covariance

- ▶ functional variable $X(t)$,
with $t \in \mathcal{T}$ and \mathcal{T} interval in \mathbb{R}
- ▶ sample $x_i(t)$, $i = 1, \dots, n$

Mean, Variance and Covariance

- ▶ functional variable $X(t)$,
with $t \in \mathcal{T}$ and \mathcal{T} interval in \mathbb{R}
- ▶ sample $x_i(t)$, $i = 1, \dots, n$
- ▶ functional mean:

$$\hat{\mu}_X(t) = \bar{x}(t) = \frac{1}{n} \sum_{i=1}^n x_i(t)$$

- ▶ functional variance:

$$\hat{\sigma}_X(t) = \frac{1}{n-1} \sum_{i=1}^n [x_i(t) - \bar{x}(t)]^2$$

Mean, Variance and Covariance

- ▶ functional variable $X(t)$,
with $t \in \mathcal{T}$ and \mathcal{T} interval in \mathbb{R}
- ▶ sample $x_i(t)$, $i = 1, \dots, n$
- ▶ functional mean:

$$\hat{\mu}_X(t) = \bar{x}(t) = \frac{1}{n} \sum_{i=1}^n x_i(t)$$

- ▶ functional variance:


$$\hat{\sigma}_X(t) = \frac{1}{n-1} \sum_{i=1}^n [x_i(t) - \bar{x}(t)]^2$$

- ▶ functional (auto-)covariance:


$$\hat{\sigma}_X(t_1, t_2) = \frac{1}{n-1} \sum_{i=1}^n [x_i(t_1) - \bar{x}(t_1)][x_i(t_2) - \bar{x}(t_2)]$$

Example for mean: Growth curves of 54 girls

growth curves with mean



centered per time-point



estimated mean:



$$\hat{\mu}_X(t) = \bar{x}(t) = \frac{1}{n} \sum_{i=1}^n x_i(t)$$

centered curves:

$$x_i^*(t) = x_i(t) - \bar{x}(t)$$

Example for variance

centered per time-point



centered curves:

$$x_i^*(t) = x_i(t) - \bar{x}(t)$$



estimated variance:

$$\hat{\sigma}_X(t) = \frac{1}{n-1} \sum_{i=1}^n [x_i(t) - \bar{x}(t)]^2$$

12/43

Example for covariance surface

$$\hat{\sigma}_X(t_1, t_2) = \frac{1}{n-1} \sum_{i=1}^n [x_i(t_1) - \bar{x}(t_1)][x_i(t_2) - \bar{x}(t_2)]$$





Functional data analysis in a nutshell

Outlook to functional data analysis

Important topics: (Ramsay and Silverman, 2005)



- Data representation → interpolation, smoothing



Outlook to functional data analysis

Important topics: (Ramsay and Silverman, 2005)


- ▶ Data representation \rightarrow interpolation, smoothing
- ▶ Visualization \rightarrow registration, outlier detection



Outlook to functional data analysis

Important topics: (Ramsay and Silverman, 2005)



- ▶ Data representation → interpolation, smoothing
- ▶ Visualization → registration, outlier detection
- ▶ Finding of patterns in the variation of the data → functional principal component analysis (FPCA)



Outlook to functional data analysis

Important topics: (Ramsay and Silverman, 2005)

- ▶ Data representation → interpolation, smoothing
- ▶ Visualization → registration, outlier detection
- ▶ Finding of patterns in the variation of the data → functional principal component analysis (FPCA)
- ▶ Classification and clustering



Outlook to functional data analysis

Important topics: (Ramsay and Silverman, 2005)

- ▶ Data representation → interpolation, smoothing
- ▶ Visualization → registration, outlier detection
- ▶ Finding of patterns in the variation of the data → functional principal component analysis (FPCA)
- ▶ Classification and clustering
- ▶ Regression → functional regression models (Morris, 2015; Greven and Scheipl, 2017)

scalar-on-function:
$$y_i = \beta_0 + \int x_i(s) \beta(s) ds + \varepsilon_i$$

function-on-scalar:
$$y_i(t) = \beta_0(t) + x_i \beta(t) + \varepsilon_i(t)$$

function-on-function:
$$y_i(t) = \beta_0(t) + \int x_i(s) \beta(s, t) ds + \varepsilon_i(t)$$

R packages

Visualization

- ▶ Shang & Hyndman (2016). *rainbow: Rainbow Plots, Bagplots and Boxplots for Functional Data*. R package version 3.4.
<https://CRAN.R-project.org/package=rainbow>

Visualization, descriptive and exploratory analysis

- ▶ Febrero-Bande & Oviedo de la Fuente (2012). *Statistical Computing in Functional Data Analysis: The R Package fda.usc*. Journal of Statistical Software, 51(4), 1–28.
- ▶ Ramsay, Wickham, Graves & Hooker (2014). *fda: Functional Data Analysis*. R package version 2.4.4. <https://CRAN.R-project.org/package=fda>

Regression

- ▶ Goldsmith, Scheipl, Huang, Wrobel, Gellar, Harezlak, McLean, Swihart, Xiao, Crainiceanu & Reiss (2016). *refund: Regression with Functional Data*. R package version 0.1-16. <https://CRAN.R-project.org/package=refund>
- ▶ Brockhaus & Rügamer (2017). *FDboost: Boosting Functional Regression models*. R package version 0.3-0.
<https://CRAN.R-project.org/package=FDboost>

see also the CRAN Task View: Functional Data Analysis


Regression with functional data

Motivation: emotion components data (I)

Data set from Gentsch et al. (2014), also used in Rügamer et al. (2016)


- ▶ Main goal: Understand how emotions evolve
- ▶ Participants played a gambling game with real money outcome
- ▶ Emotions “measured” via EMG (muscle activity in the face)
- ▶ Influencing factor *appraisals* measured via EEG (brain activity)
- ▶ Different game situation, a lot of trials

Motivation: emotion components data (II)



(Brockhaus et al., 2017)


Motivation: emotion components data (II)



(Brockhaus et al., 2017)

Function-on-function-regression


Motivation: emotion components data (II)



(Brockhaus et al., 2017)

Function-on-function-regression ... what for?

Motivation: emotion components data (II)




(Brockhaus et al., 2017)

Function-on-function-regression ... what for?

- The absolute value is not really of interest

Motivation: emotion components data (II)



(Brockhaus et al., 2017)

Function-on-function-regression ... what for?

- ▶ The absolute value is not really of interest
- ▶ Describe the course of each curve, more specifically their relationship → average course of function

Generic additive regression model

- ▶ functional response $Y(t)$, $t \in \mathcal{T} = [T_1, T_2]$
- ▶ vector of covariates \mathbf{x} containing functional covariates $x(s)$ and scalar covariates z

Generic model

$$\mathbb{E}(Y(t) \mid \mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$$

- ▶ $h(\mathbf{x}, t)$ linear predictor which is the sum of partial effects $h_j(\mathbf{x}, t)$
- ▶ each $h_j(\mathbf{x}, t)$ is a real valued function over \mathcal{T} and can depend on or several covariates

Generic additive regression model

- ▶ functional response $Y(t)$, $t \in \mathcal{T} = [T_1, T_2]$
- ▶ vector of covariates \mathbf{x} containing functional covariates $x(s)$ and scalar covariates z

Generic model

$$\mathbb{E}(Y(t) \mid \mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$$

- ▶ $h(\mathbf{x}, t)$ linear predictor which is the sum of partial effects $h_j(\mathbf{x}, t)$
- ▶ each $h_j(\mathbf{x}, t)$ is a real valued function over \mathcal{T} and can depend on or several covariates

→ Scalar response as degenerated case with $\mathcal{T} = [T_1, T_1]$

Generic additive regression model

- ▶ functional response $Y(t)$, $t \in \mathcal{T} = [T_1, T_2]$
- ▶ vector of covariates \mathbf{x} containing functional covariates $x(s)$ and scalar covariates z

Generic model

$$\mathbb{E}(Y(t) \mid \mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$$

- ▶ $h(\mathbf{x}, t)$ linear predictor which is the sum of partial effects $h_j(\mathbf{x}, t)$
- ▶ each $h_j(\mathbf{x}, t)$ is a real valued function over \mathcal{T} and can depend on or several covariates

→ Scalar response as degenerated case with $\mathcal{T} = [T_1, T_1]$

Partial effects $h_j(x, t)$ of scalar covariates

- ▶ smooth intercept $\beta_0(t)$
- ▶ group-specific smooth intercepts $\beta_{0a}(t)$
- ▶ smooth linear effect of scalar covariate $z\beta(t)$
- ▶ smooth non-linear effect of scalar covariate $g(z, t)$
- ▶ interactions, e.g., $z_1 z_2 \beta(t)$ and $g(z_1, z_2, t)$

(Scheipl et al., 2015; Brockhaus et al., 2015)

Partial effects $h_j(x, t)$ of functional covariates

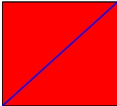
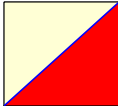
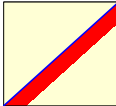
- ▶ concurrent effect $x(t)\beta(t)$
- ▶ linear effect of functional covariate $\int_{\mathcal{S}} x(s)\beta(s, t) ds$

Partial effects $h_j(x, t)$ of functional covariates

- ▶ concurrent effect $x(t)\beta(t)$
- ▶ linear effect of functional covariate $\int_{\mathcal{S}} x(s)\beta(s, t) ds$
- ▶ constrained effect of functional covariate $\int_{l(t)}^{u(t)} x(s)\beta(s, t) ds$,
with integration limits $[l(t), u(t)]$

Partial effects $h_j(x, t)$ of functional covariates

- ▶ concurrent effect $x(t)\beta(t)$
- ▶ linear effect of functional covariate $\int_{\mathcal{S}} x(s)\beta(s, t) ds$
- ▶ constrained effect of functional covariate $\int_{l(t)}^{u(t)} x(s)\beta(s, t) ds$, with integration limits $[l(t), u(t)]$

effect	linear	historical	lag
$[l(t), u(t)]$	$[T_1, T_2]$	$[T_1, t]$	$[t - \delta, t]$
s			
	t		

(Scheipl et al., 2015; Brockhaus et al., 2016b,a)

Interactions of functional and scalar covariates

- ▶ linear interaction of scalar and functional covariate

$$z \int_{l(t)}^{u(t)} x(s) \beta(s, t) ds$$

- ▶ group-specific functional effects

$$I(z = a) \cdot \int_{l(t)}^{u(t)} x(s) \beta_a(s, t) ds$$

with indicator function $I(\cdot)$

Interactions of functional and scalar covariates

- ▶ linear interaction of scalar and functional covariate

$$z \int_{l(t)}^{u(t)} x(s) \beta(s, t) ds$$

- ▶ group-specific functional effects

$$I(z = a) \cdot \int_{l(t)}^{u(t)} x(s) \beta_a(s, t) ds$$

with indicator function $I(\cdot)$

- For all the listed effects ensure identifiability by suitable constraints

Specification of partial effects

The generic model: $\mathbb{E}(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

Row tensor product basis

$$h_j(\mathbf{x}, t) = \left\{ \mathbf{b}_j(\mathbf{x}, t)^\top \odot \mathbf{b}_Y(t)^\top \right\} \boldsymbol{\theta}_j$$

- ▶ \mathbf{b}_j / \mathbf{b}_Y vector of κ_j/κ_Y basis functions in covariates / over \mathcal{T}
- ▶ \odot row-wise tensor product ('Kronecker product on rows')
- ▶ $\boldsymbol{\theta}_j$ coefficient vector
- ▶ Ridge-type penalty with penalty term $\boldsymbol{\theta}_j^\top \mathbf{P}_{jY} \boldsymbol{\theta}_j$ for regularization in both directions

Specification of partial effects

The generic model: $\mathbb{E}(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

Row tensor product basis

$$h_j(\mathbf{x}, t) = \left\{ \mathbf{b}_j(\mathbf{x}, t)^\top \odot \mathbf{b}_Y(t)^\top \right\} \theta_j$$

- ▶ \mathbf{b}_j / \mathbf{b}_Y vector of κ_j / κ_Y basis functions in covariates / over \mathcal{T}
- ▶ \odot row-wise tensor product ('Kronecker product on rows')
- ▶ θ_j coefficient vector
- ▶ Ridge-type penalty with penalty term $\theta_j^\top \mathbf{P}_{jY} \theta_j$ for regularization in both directions
- ▶ if possible, representation as **generalized linear array model** (Currie et al., 2006).

Estimation

How do we estimate such models?

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
- ▶ Problem: Computational feasibility

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
- ▶ Problem: Computational feasibility
 - For example, consider a model with

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
- ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
- ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$
 - ▶ factor with $\kappa_z = 10$ levels

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
- ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$
 - ▶ factor with $\kappa_z = 10$ levels
 - ▶ $\kappa_s = \kappa_t = 20$ B-spline bases for $\beta_a(s, t)$ smoothness in s and t

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
 - ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$
 - ▶ factor with $\kappa_z = 10$ levels
 - ▶ $\kappa_s = \kappa_t = 20$ B-spline bases for $\beta_a(s, t)$ smoothness in s and t
- $\Rightarrow \text{ncol}(\mathbf{X}) = \kappa_z \cdot \kappa_s \cdot \kappa_t$

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
 - ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$
 - ▶ factor with $\kappa_z = 10$ levels
 - ▶ $\kappa_s = \kappa_t = 20$ B-spline bases for $\beta_a(s, t)$ smoothness in s and t
- $\Rightarrow \text{ncol}(\mathbf{X}) = \kappa_z \cdot \kappa_s \cdot \kappa_t = 10 \cdot 20 \cdot 20 = 4000$

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
 - ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$
 - ▶ factor with $\kappa_z = 10$ levels
 - ▶ $\kappa_s = \kappa_t = 20$ B-spline bases for $\beta_a(s, t)$ smoothness in s and t
- $\Rightarrow \text{ncol}(\mathbf{X}) = \kappa_z \cdot \kappa_s \cdot \kappa_t = 10 \cdot 20 \cdot 20 = 4000$
- ▶ So how can we handle and fit multiple of such partial effects at the same time?

Estimation

How do we estimate such models?

- ▶ As per usual: Write down the (penalized) log-likelihood and maximize it (see, e.g., Scheipl et al., 2016)
- ▶ Problem: Computational feasibility
 - For example, consider a model with
 - ▶ a factor-specific historical effect $\int_0^t x(s) \beta_a(s, t) ds$
 - ▶ factor with $\kappa_z = 10$ levels
 - ▶ $\kappa_s = \kappa_t = 20$ B-spline bases for $\beta_a(s, t)$ smoothness in s and t
 - $\Rightarrow \text{ncol}(\mathbf{X}) = \kappa_z \cdot \kappa_s \cdot \kappa_t = 10 \cdot 20 \cdot 20 = 4000$
- ▶ So how can we handle and fit multiple of such partial effects at the same time?

→ component-wise boosting

Estimation by component-wise gradient boosting

Idea:

Estimation by component-wise gradient boosting

Idea:

- ▶ iteratively **boost** the model performance
($\hat{\equiv}$ reduce expected L_2 -loss)

Estimation by component-wise gradient boosting

Idea:

- ▶ iteratively **boost** the model performance
($\hat{=}$ reduce expected L_2 -loss)
- ▶ by fitting and evaluating the partial effects component-wise

Estimation by component-wise gradient boosting

Idea:

- ▶ iteratively **boost** the model performance
($\hat{=}$ reduce expected L_2 -loss)
- ▶ by fitting and evaluating the partial effects component-wise
- ▶ using one partial effect at a time to update the model

Estimation by component-wise gradient boosting

Idea:

- ▶ iteratively **boost** the model performance
($\hat{=}$ reduce expected L_2 -loss)
- ▶ by fitting and evaluating the partial effects component-wise
- ▶ using one partial effect at a time to update the model


→ Results in component-wise gradient descent steps

Estimation by component-wise gradient boosting

Idea:

- ▶ iteratively **boost** the model performance
($\hat{=}$ reduce expected L_2 -loss)
- ▶ by fitting and evaluating the partial effects component-wise
- ▶ using one partial effect at a time to update the model

→ Results in component-wise gradient descent steps



Source: <https://github.com/joshdk/pygradesc>

Component-wise gradient boosting: algorithm

Goal of boosting: Minimize the expected loss

Use the (penalized) regression models for effects h_j as base-learners

Algorithm For boosting iterations $m = 1, \dots, m_{\text{stop}}$:

- ▶ **compute** the negative gradient $u_i(t)$ of the expected loss using the current estimate of the linear predictor $\hat{h}^{[m]}(x_i, t)$
- ▶ **fit** each base-learner to $u_i(t)$
- ▶ **select** the best fitting base-learner
- ▶ **update** the according parameters using a fixed step-length $\nu \in (0, 1]$

The final model is a linear combination of base-learner fits.

(Brockhaus et al., 2015)

Remember the generic model

- ▶ functional response $Y(t)$, $t \in \mathcal{T} \subset \mathbb{R}$
- ▶ vector of covariates \mathbf{x} containing functional covariates $x(s)$ and scalar covariates z

Generic model

$$\mathbb{E}(Y(t) \mid \mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$$

- ▶ $h(\mathbf{x}, t)$ linear predictor which is the sum of partial effects $h_j(\mathbf{x}, t)$
- ▶ each $h_j(\mathbf{x}, t)$ is a real valued function over \mathcal{T} and can depend on or several covariates

Remember the generic model

- ▶ functional response $Y(t)$, $t \in \mathcal{T} \subset \mathbb{R}$
- ▶ vector of covariates \mathbf{x} containing functional covariates $x(s)$ and scalar covariates z

Generic model

$$\xi(Y(t) | \mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$$

- ▶ $h(\mathbf{x}, t)$ linear predictor which is the sum of partial effects $h_j(\mathbf{x}, t)$
- ▶ each $h_j(\mathbf{x}, t)$ is a real valued function over \mathcal{T} and can depend on or several covariates

Remember the generic model

- ▶ functional response $Y(t)$, $t \in \mathcal{T} \subset \mathbb{R}$
- ▶ vector of covariates \mathbf{x} containing functional covariates $x(s)$ and scalar covariates z

Generic model

$$\xi(Y(t) | \mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$$

- ▶ $h(\mathbf{x}, t)$ linear predictor which is the sum of partial effects $h_j(\mathbf{x}, t)$
- ▶ each $h_j(\mathbf{x}, t)$ is a real valued function over \mathcal{T} and can depend on or several covariates
- ▶ ξ is some transformation function, e.g., a quantile or \mathbb{E}

Transformation and loss functions

The generic model: $\xi(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

Transformation and loss functions

The generic model: $\xi(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

model	ξ	loss function ρ
LM	\mathbb{E}	L_2 -loss
GLM	$g \circ \mathbb{E}$	negative log-likelihood
median regression	$Q_{0.5}$	L_1 -loss
quantile regression	Q_τ	check function

Transformation and loss functions

The generic model: $\xi(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

model	ξ	loss function ρ
LM	\mathbb{E}	L_2 -loss
GLM	$g \circ \mathbb{E}$	negative log-likelihood
median regression	$Q_{0.5}$	L_1 -loss
quantile regression	Q_τ	check function

Transformation and loss functions

The generic model: $\xi(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

model	ξ	loss function ρ
LM	\mathbb{E}	L_2 -loss
GLM	$g \circ \mathbb{E}$	negative log-likelihood
median regression	$Q_{0.5}$	L_1 -loss
quantile regression	Q_τ	check function

→ GAMLSS (Stasinopoulos et al., 2016) also possible

Transformation and loss functions

The generic model: $\xi(Y(t)|\mathbf{x}) = h(\mathbf{x}, t) = \sum_j h_j(\mathbf{x}, t)$

model	ξ	loss function ρ
LM	\mathbb{E}	L_2 -loss
GLM	$g \circ \mathbb{E}$	negative log-likelihood
median regression	$Q_{0.5}$	L_1 -loss
quantile regression	Q_τ	check function

- GAMLSS (Stasinopoulos et al., 2016) also possible
- Loss function for trajectories $\hat{=}$ integrated loss over domain of response:

$$\ell(Y, h(\mathbf{x})) = \int_{\mathcal{T}} \underbrace{\rho(Y(t), h(\mathbf{x}, t))}_{\text{pointwise loss for } t} dt$$

Tuning, early stopping and model selection

- ▶ Tuning
 - ▶ We fix the ν and degrees of freedom for each baselearner
 - ▶ number of boosting iterations determines model complexity
 - ▶ optimal stopping iteration is determined by resampling methods (on the level of curves)

Tuning, early stopping and model selection

- ▶ Tuning
 - ▶ We fix the ν and degrees of freedom for each baselearner
 - ▶ number of boosting iterations determines model complexity
 - ▶ optimal stopping iteration is determined by resampling methods (on the level of curves)
- ▶ Early stopping
 - ▶ Regularization technique to avoid overfitting
 - ▶ Induces parameter shrinkage
 - ▶ and model selection

Tuning, early stopping and model selection

- ▶ Tuning
 - ▶ We fix the ν and degrees of freedom for each baselearner
 - ▶ number of boosting iterations determines model complexity
 - ▶ optimal stopping iteration is determined by resampling methods (on the level of curves)
- ▶ Early stopping
 - ▶ Regularization technique to avoid overfitting
 - ▶ Induces parameter shrinkage
 - ▶ and model selection
- ▶ Alternatively: Model selection via stability selection (Meinshausen and Bühlmann, 2010; Shah and Samworth, 2013)

Implementation

Implemented in

- ▶ R package **FDboost** (Brockhaus et al., 2017)
- ▶ based on R package `mboost` (Hothorn et al., 2016)
- ▶ Extension: `FDboostLSS` for functional `gamboostLSS` (Hofner et al., 2015)

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

Implementation in FDbboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'
 - ▶ We just have to implement new base-learners and

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'
 - ▶ We just have to implement new base-learners and
 - ▶ allow for different data input (e.g. matrices for curve observations)

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'
 - ▶ We just have to implement new base-learners and
 - ▶ allow for different data input (e.g. matrices for curve observations)
- ▶ Function-on-function regression:

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'
 - ▶ We just have to implement new base-learners and
 - ▶ allow for different data input (e.g. matrices for curve observations)
- ▶ Function-on-function regression:
 - ▶ Implement base-learner, which also vary in the direction of t

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'
 - ▶ We just have to implement new base-learners and
 - ▶ allow for different data input (e.g. matrices for curve observations)
- ▶ Function-on-function regression:
 - ▶ Implement base-learner, which also vary in the direction of t
 - ▶ The loss function is now an integral

Implementation in FDboost (I)

Goal: Make use of the modular implementation of `mboost`

- ▶ Scalar-on-function regression:
the loss and empirical risk is just as for 'scalar-on-scalar regression'
 - ▶ We just have to implement new base-learners and
 - ▶ allow for different data input (e.g. matrices for curve observations)
- ▶ Function-on-function regression:
 - ▶ Implement base-learner, which also vary in the direction of t
 - ▶ The loss function is now an integral
 - Numerical integration scheme to approximate expected loss

Implementation in FDboost (II)

- Main fitting function:

`FDboost(formula, timeformula, data, ...)`

- `timeformula`

= NULL for scalar-on-function regression,

= \sim `bbs(t)` for function-on-function regression

- Some of the base-learners for functional data:

$z\beta(t)$ `bolsc(z)`

$f(z, t)$ `bbsc(z)`

$z_1 z_2 \beta(t)$ `bols(z1) %Xc% bols(z2)`

$\int_S x(s)\beta(s, t)ds$ `bsignal(x, s = s)`

$x(t)\beta(t)$ `bconcurrent(x, s = s, time = t)`

$\int_{l(t)}^{u(t)} x(s)\beta(s, t)ds$ `bhist(x, s = s, time = t, limits = ...)`

Implementation in FDboost (II)

- Main fitting function:

`FDboost(formula, timeformula, data, ...)`

- `timeformula`

= `NULL` for scalar-on-function regression,

= `~ bbs(t)` for function-on-function regression

- Some of the base-learners for functional data:

$z\beta(t)$ `bolsc(z) %0% bbs(t)`

$f(z, t)$ `bbsc(z) %0% bbs(t)`

$z_1 z_2 \beta(t)$ `bols(z1) %Xc% bols(z2) %0% bbs(t)`

$\int_S x(s) \beta(s, t) ds$ `bsignal(x, s = s) %0% bbs(t)`

$x(t) \beta(t)$ `bconcurrent(x, s = s, time = t)`

$\int_{l(t)}^{u(t)} x(s) \beta(s, t) ds$ `bhist(x, s = s, time = t, limits = ...)`

Example: FDboost call

```
FDboost(EMG ~ 1 +  
        brandomc(id, df = 5) +  
        bhistx(EEG, df = 20),  
        timeformula = ~ bbs(t, df = 4),  
        control = boost_control(mstop = 5000,  
                                trace = TRUE),  
        data = data)
```

Case studies


Functional response

Emotion components data

Goal: Try to explain

- ▶ facial expressions (measured with EMG)
- ▶ by brain activity (measured with EEG)

→ Function-on-function-regression



(Brockhaus et al., 2017)

Emotion components data

Model equation:

$$y_{\text{EMG}}(t) = \beta_0(t) + x_{\text{EEG}}(t)\beta_1(t) + \varepsilon(t)$$

- One-to-one relation between EEG and EMG → Concurrent effect

Emotion components data

Model equation:

$$y_{\text{EMG}}(t) = \beta_0(t) + \int x_{\text{EEG}}(s) \beta_1(s, t) ds + \varepsilon(t)$$

- ▶ One-to-one relation between EEG and EMG → Concurrent effect
- ▶ Response time specific effect → Linear functional effect

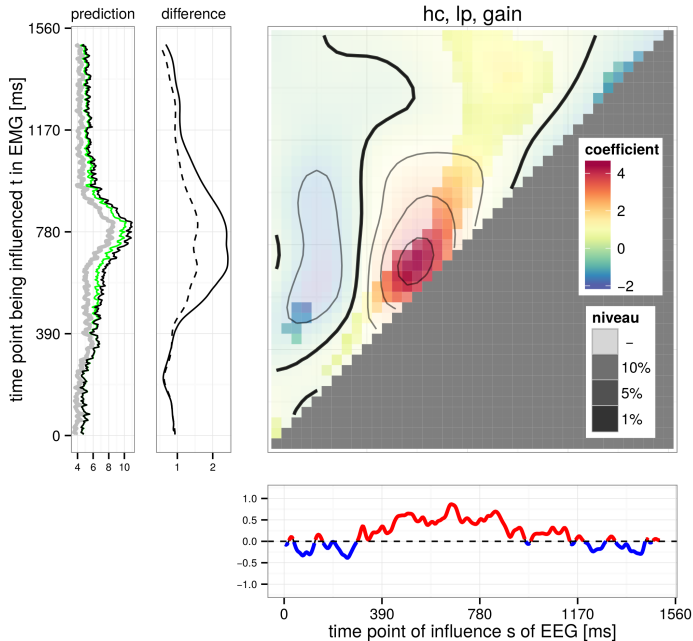
Emotion components data

Model equation:

$$y_{\text{EMG}}(t) = \beta_0(t) + \int_0^{t-\delta} x_{\text{EEG}}(s) \beta_1(s, t) ds + \varepsilon(t)$$

- ▶ One-to-one relation between EEG and EMG → Concurrent effect
- ▶ Response time specific effect → Linear functional effect
- ▶ EMG can only be influenced by EEG activities in the past → Historical effect

Results




Case studies

Scalar response and functional covariates

Spectral data of fossil fuels

Goal: predict heat value y using the spectral measurements of NIR and UV spectra



(Brockhaus et al., 2015)

Data of fossil fuel: model

Model equation:


$$y = \beta_0 + f(z_{\text{H}_2\text{O}}) + \int_{\mathcal{S}_{\text{NIR}}} x_{\text{NIR}}(s_{\text{NIR}}) \beta_{\text{NIR}}(s_{\text{NIR}}) ds_{\text{NIR}} + \int_{\mathcal{S}_{\text{UV}}} x_{\text{UV}}(s_{\text{UV}}) \beta_{\text{UV}}(s_{\text{UV}}) ds_{\text{UV}} + \varepsilon,$$

- ▶ heat value y
- ▶ non-linear effect of water content (H₂O)
- ▶ linear functional effect of NIR and UV spectrum


Data of fossil fuel: results

Estimated effects with stopping iteration chosen by 10 fold bootstrap


bbs(h2o)



bsignal(UVVIS)




bsignal(NIR)




Data of fossil fuel: results

- ▶ estimated effects on 100 bootstrap samples (gray lines)
- ▶ point-wise median (black lines)
- ▶ point-wise 5 and 95% quantiles (dashed red lines)


bbs(h2o)



bsignal(UVVIS)



bsignal(NIR)



Summary and discussion

Summary and discussion (I)

What is FDA?

- ▶ Measurement units are functions, i.e., curves, surfaces, trajectories,...
- ▶ Smooth data generating process
- ▶ Many observations of the same data generating process
- ▶ Mean, variance and covariance for functional data
- ▶ Functional counterparts for many methods from multivariate statistics

Summary and discussion (II)

Estimating functional regression models with FDboost

- ▶ LM, GLM, GAMLSS and quantile regression included
- ▶ variety of covariate effects,
 - ▶ (non-)linear effects of scalar covariates
 - ▶ linear effects of functional covariates, historical effects
 - ▶ interaction effects

Summary and discussion (II)

Estimating functional regression models with FDboost

- ▶ LM, GLM, GAMLSS and quantile regression included
- ▶ variety of covariate effects,
 - ▶ (non-)linear effects of scalar covariates
 - ▶ linear effects of functional covariates, historical effects
 - ▶ interaction effects
- ▶ estimation by component-wise gradient boosting
 - ▶ high dimensional data settings
 - ▶ data-driven variable selection
 - ▶ shrinkage of effects
- ▶ comprehensive implementation in R add-on package FDboost

References I

- S. Brockhaus, D. Rügamer, and S. Greven. Boosting Functional Regression Models with FDboost. *ArXiv e-prints*, May 2017.
- Sarah Brockhaus, Fabian Scheipl, Torsten Hothorn, and Sonja Greven. The functional linear array model. *Statistical Modelling*, 15(3):279–300, 2015.
- Sarah Brockhaus, Andreas Fuest, Andreas Mayr, and Sonja Greven. Signal regression models for location, scale and shape with an application to stock returns. arXiv preprint, arXiv:1605.04281, 2016a.
- Sarah Brockhaus, Michael Melcher, Friedrich Leisch, and Sonja Greven. Boosting flexible functional regression models with a high number of functional historical effects. *Statistics and Computing*, 2016b. Accepted, DOI: <http://dx.doi.org/10.1007/s11222-016-9662-1>.
- I. D. Currie, M. Durban, and P. H. C. Eilers. Generalized linear array models with applications to multidimensional smoothing. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 68(2):259–280, 2006.
- Kornelia Gentsch, Didier Grandjean, and Klaus R. Scherer. Coherence explored between emotion components: Evidence from event-related potentials and facial electromyography. *Biological Psychology*, 98(0):70 – 81, 2014.
- Sonja Greven and Fabian Scheipl. A general framework for functional regression modelling. *Statistical Modelling*, 17(1-2):1–35, 2017.

References II

- Benjamin Hofner, Andreas Mayr, Nora Fenske, and Matthias Schmid. *gamboostLSS: Boosting Methods for GAMLSS Models*, 2015. R package version 1.2-0, Available at <http://CRAN.R-project.org/package=gamboostLSS>.
- Torsten Hothorn, Peter Bühlmann, Thomas Kneib, Matthias Schmid, and Benjamin Hofner. *mboost: Model-Based Boosting*, 2016. R package version 2.6-0, Available at <http://CRAN.R-project.org/package=mboost>.
- Nicolai Meinshausen and Peter Bühlmann. Stability selection (with discussion). *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(4):417–473, 2010.
- Jeffrey S. Morris. Functional regression. *Annual Review of Statistics and Its Application*, 2(1):321–359, 2015.
- James O Ramsay and Bernard W Silverman. *Functional Data Analysis*. Springer, New York, 2005.
- David Rügamer, Sarah Brockhaus, Kornelia Gentsch, Klaus Scherer, and Sonja Greven. Detecting synchronisation in EEG- and EMG-signals via boosted functional historical models. Unpublished manuscript, 2016.
- Fabian Scheipl, Ana-Maria Staicu, and Sonja Greven. Functional additive mixed models. *Journal of Computational and Graphical Statistics*, 24(2): 477–501, 2015.

References III

- Fabian Scheipl, Jan Gertheiss, and Sonja Greven. Generalized functional additive mixed models. *Electronic Journal of Statistics*, 10(1):1455–1492, 2016.
- Rajen D Shah and Richard J Samworth. Variable selection with error control: another look at stability selection. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 75(1):55–80, 2013.
- D. M. Stasinopoulos, R. A. Rigby, V. Voudouris, C. Akantziliotou, and M. Enea. *gamlss: Generalised Additive Models for Location Scale and Shape*, 2016. R package version 4.3-8, Available at <http://CRAN.R-project.org/package=gamlss>.